

A dark blue, irregular ink splash or blotch serves as the background for the text. The splash has a textured, watercolor-like appearance with some lighter blue and white areas around the edges, suggesting ink spreading on a light surface. The text is centered within the darkest part of the splash.

System Design Basics

Agenda

- (Some) Key Characteristics of Distributed Systems
- Load Balancing
- Caching
- Asynchronism
- Scaling the DB
- Consistent Hashing

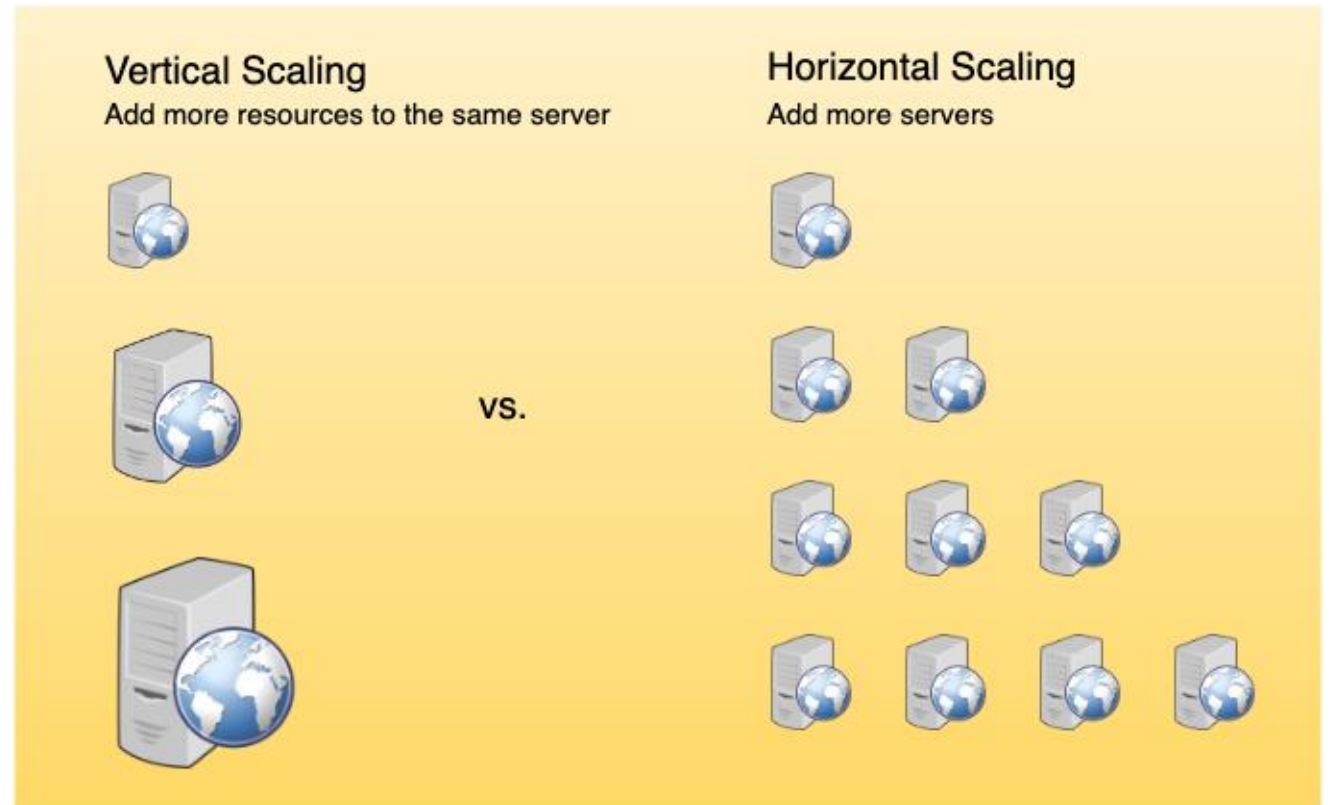
(Some) Key Characteristics of Distributed Systems

What are DS and why do we need them?




Scalability

- What? - the capability of a system, process, or a network to grow and manage increased demand without performance loss
- Why? - a system may have to scale because of many reasons like increased data volume or increased amount of work, e.g., number of transaction
- How? - **Horizontally** and **Vertically**





Reliability

- What? - the probability a system will fail in a given period
- Why? - a system needs to be reliable in order to keep delivering its services even when one or several of its software or hardware components fail
- How? - **Redundancy** of both the **software components** and **data**



Ex: Server Redundancy



Availability

- What? - the time a system remains operational to perform its required function in a specific period under normal conditions
- How? -
 redundancy and **replication**



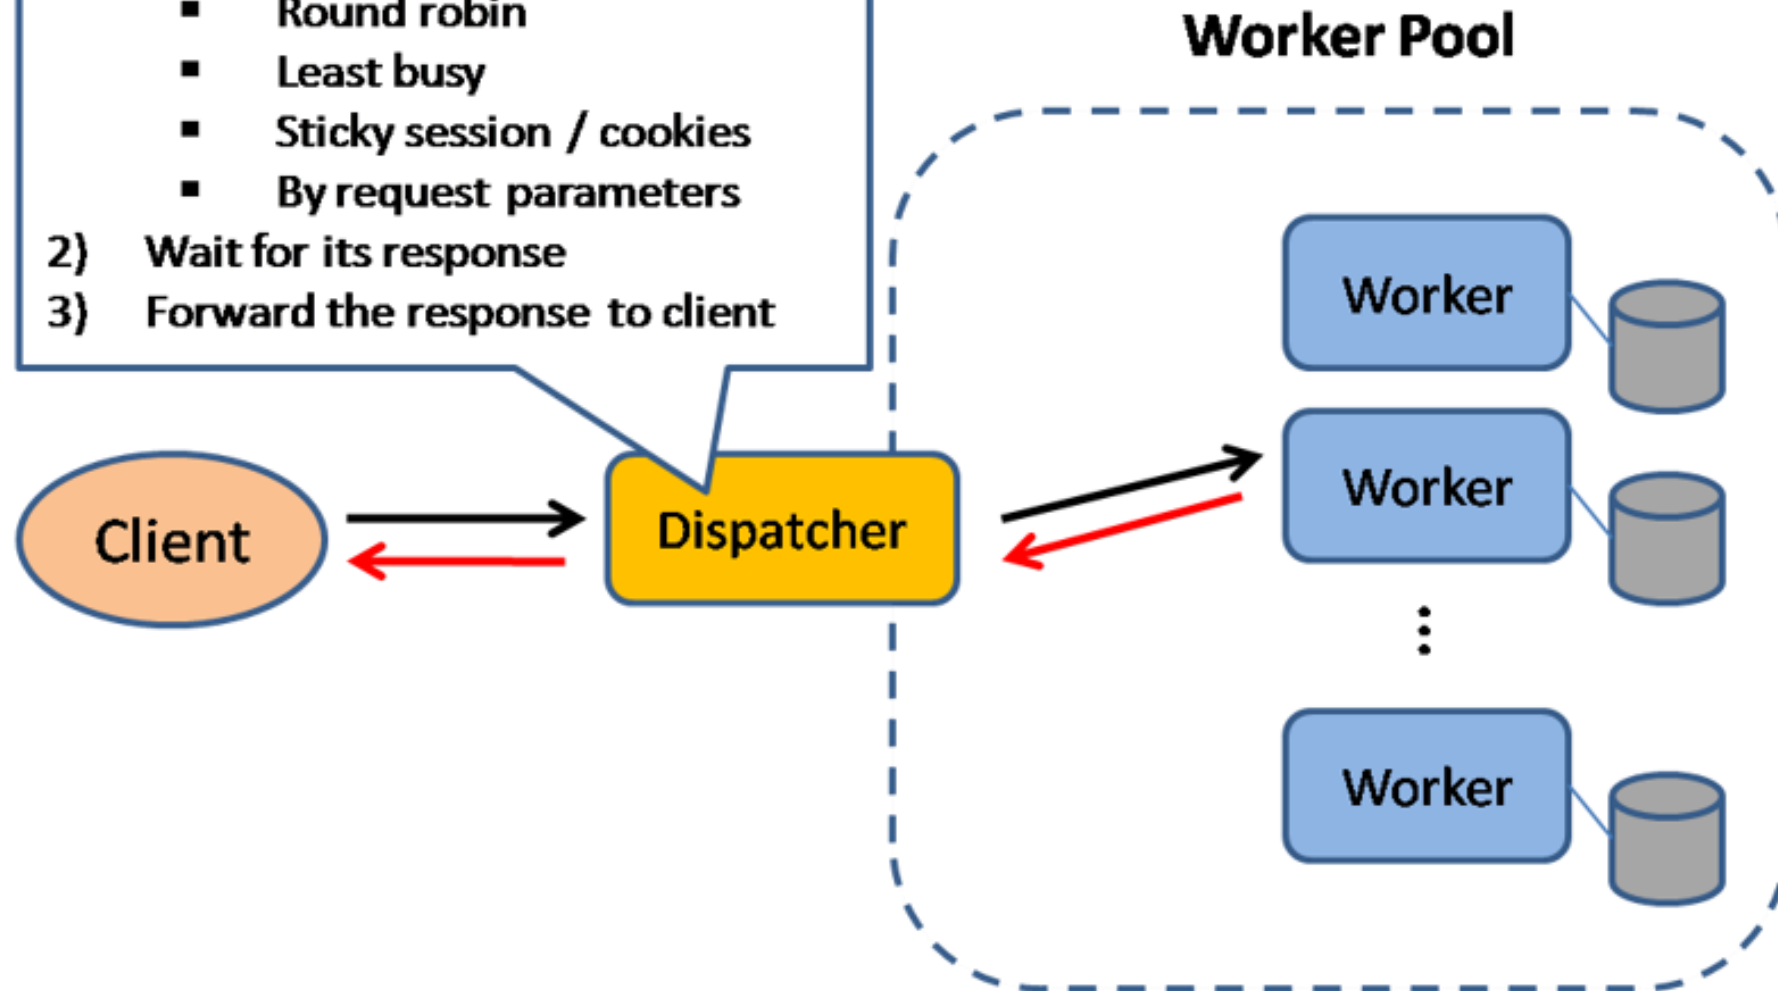
Load Balancer

1) Pick a worker to forward request


- Random
- Round robin
- Least busy
- Sticky session / cookies
- By request parameters

2) Wait for its response

3) Forward the response to client







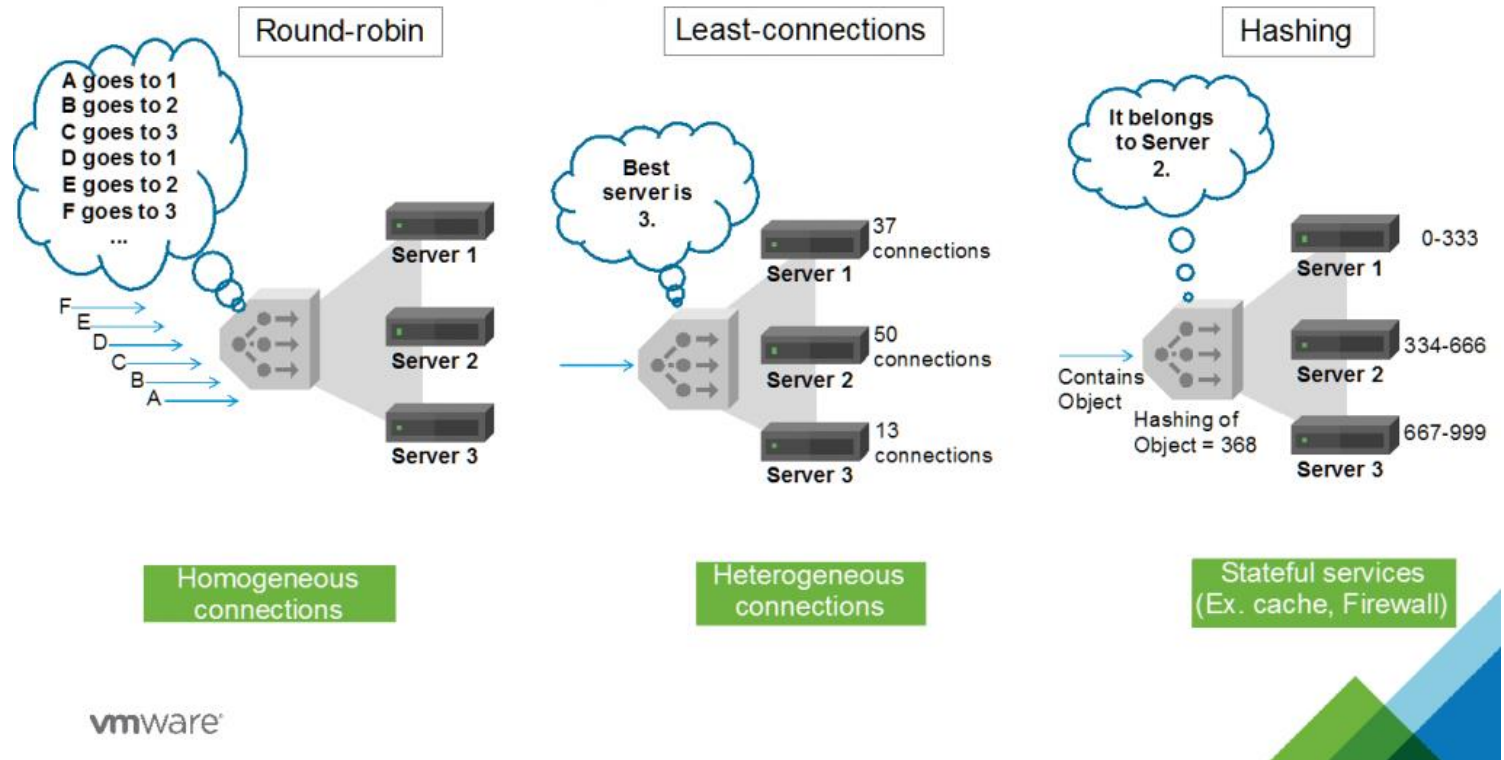
Where?

- Between the user and the web server
- Between web servers and an internal platform layer, like application servers or cache servers
- Between internal platform layer and database.


Load Balancing Algorithms

- Least Connection Method
- Least Response Time Method
- Least Bandwidth Method
- (Weighted) Round Robin Method
- IP Hash
- Consistent Hashing

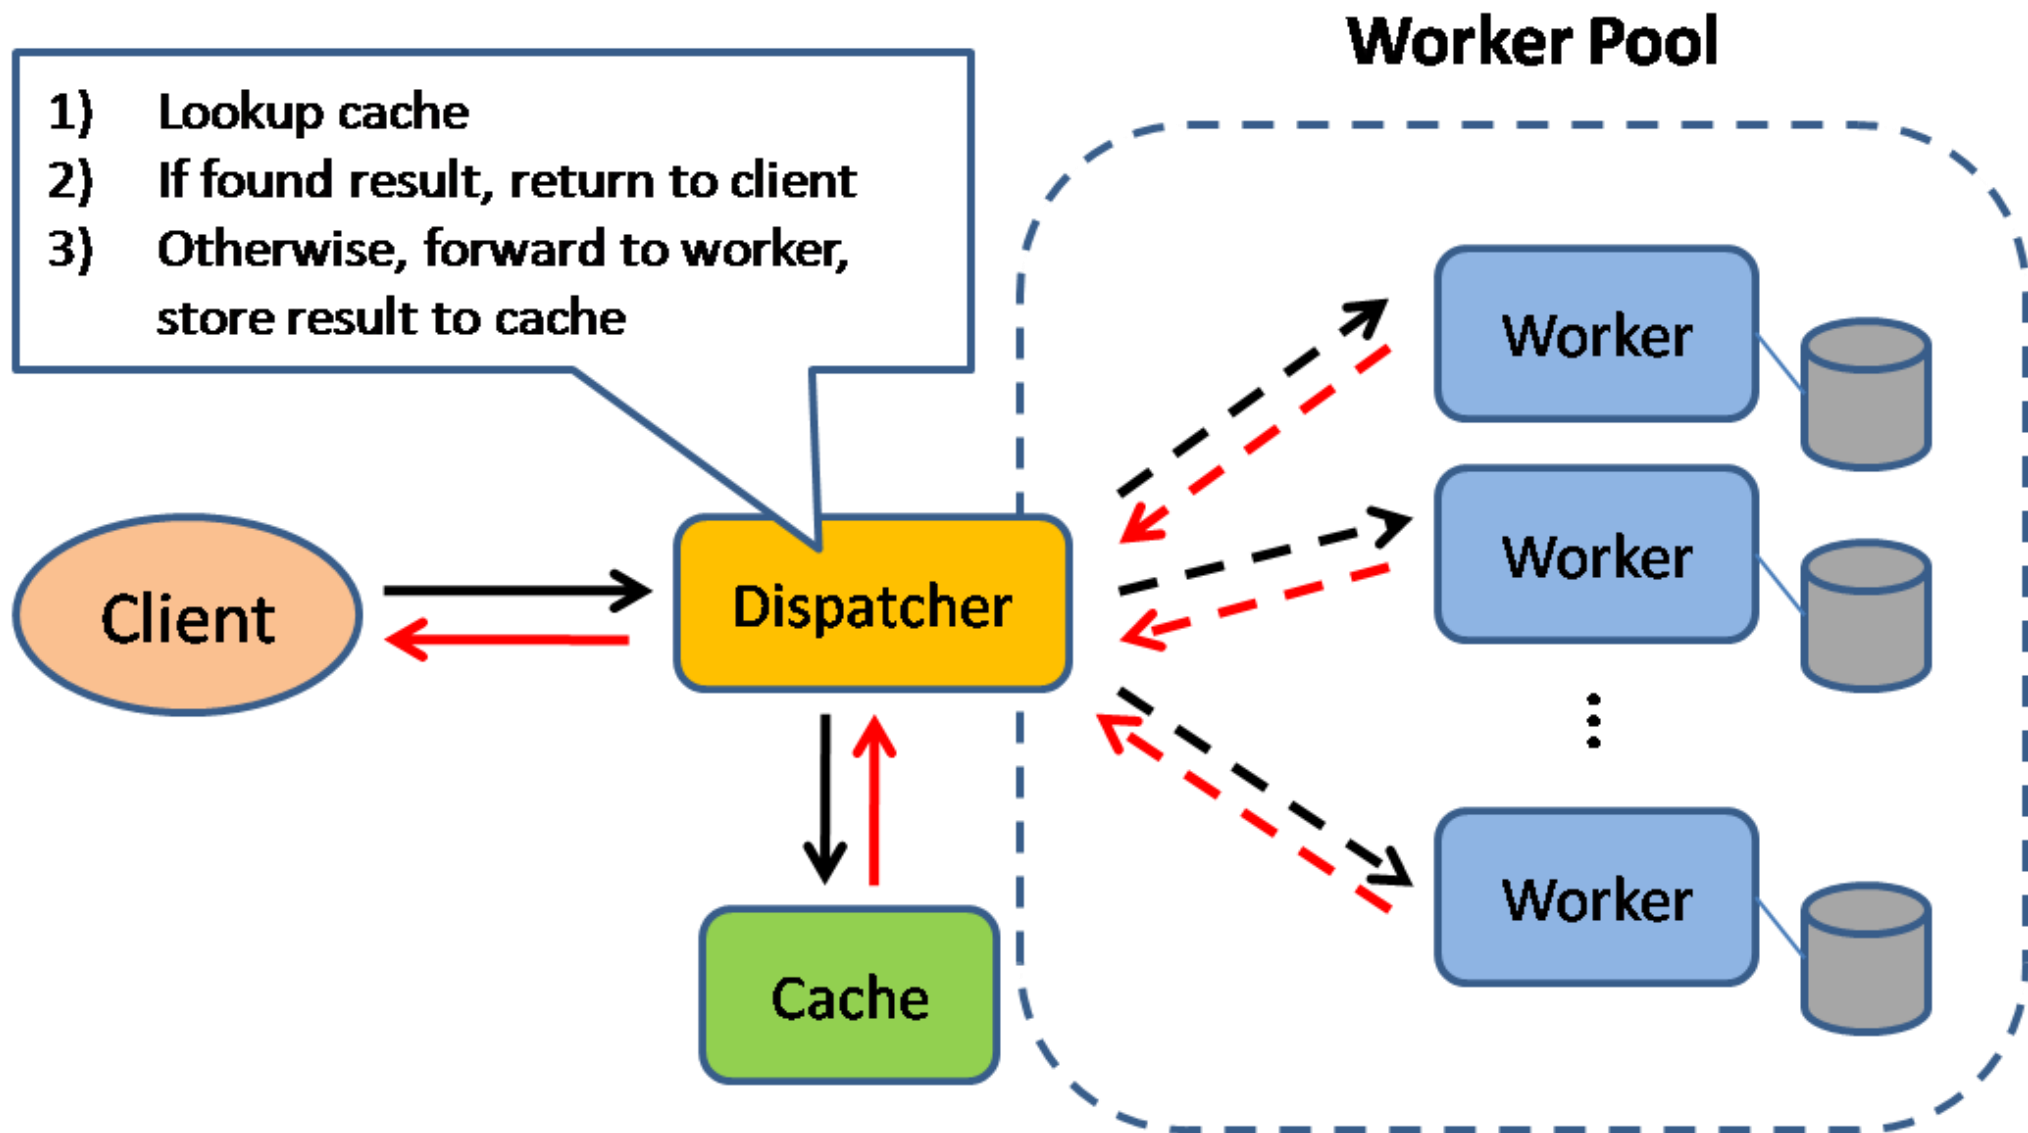
Basic Load Balancing Algorithms



Redundant Load Balancers



Caching

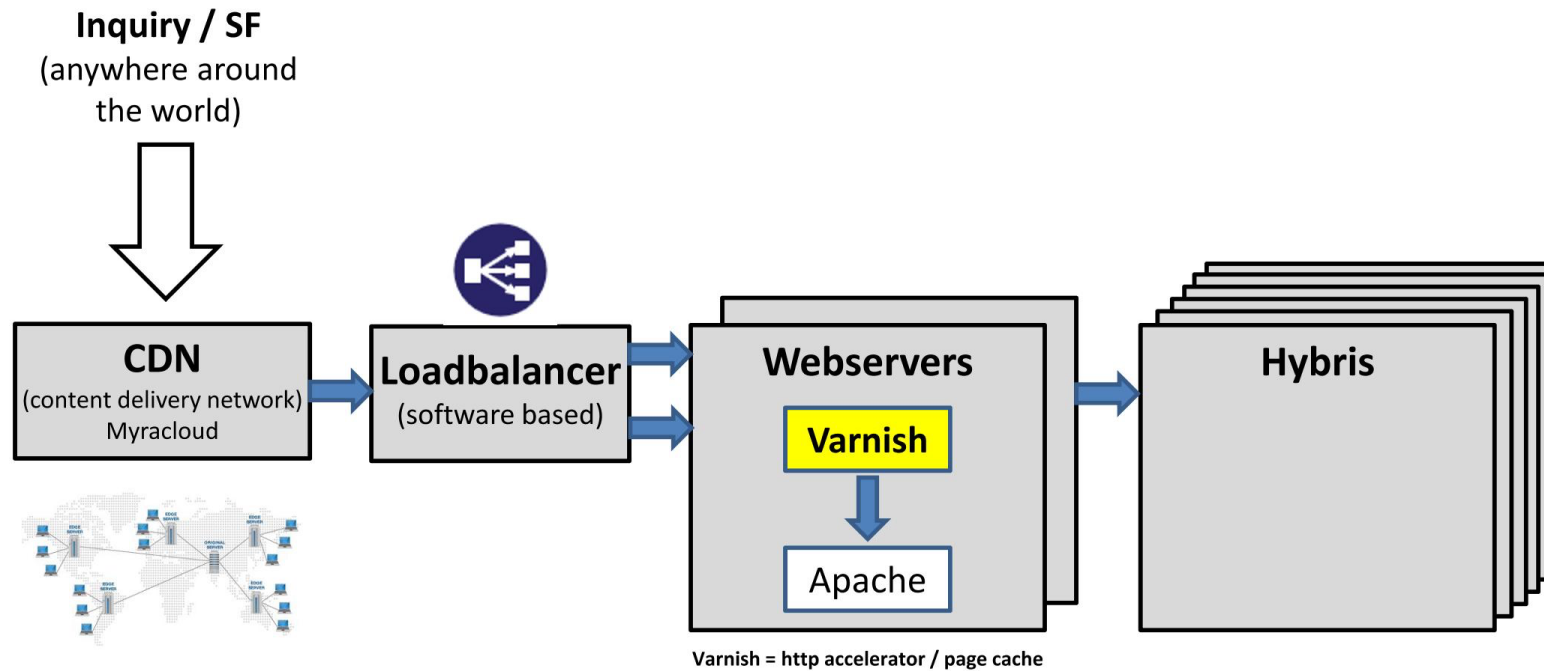




Where?

- CDN caching (Myra)
- Web Server caching – Reverse Proxies (Varnish)
- Application Server caching (Redis)
- Database Caching

Varnish in its ECI environment



can't have every user
who loads her timeline
look up all their followers
and then their photos

instead, everyone gets
their own list in Redis

Redis is awesome for
this; rapid insert, rapid
subsets

media ID is pushed onto
a list for every person
who's following this user

when time to render a
feed, we take small # of
IDs, go look up info in
memcached

Instagram Feed and Redis

Cache eviction policies

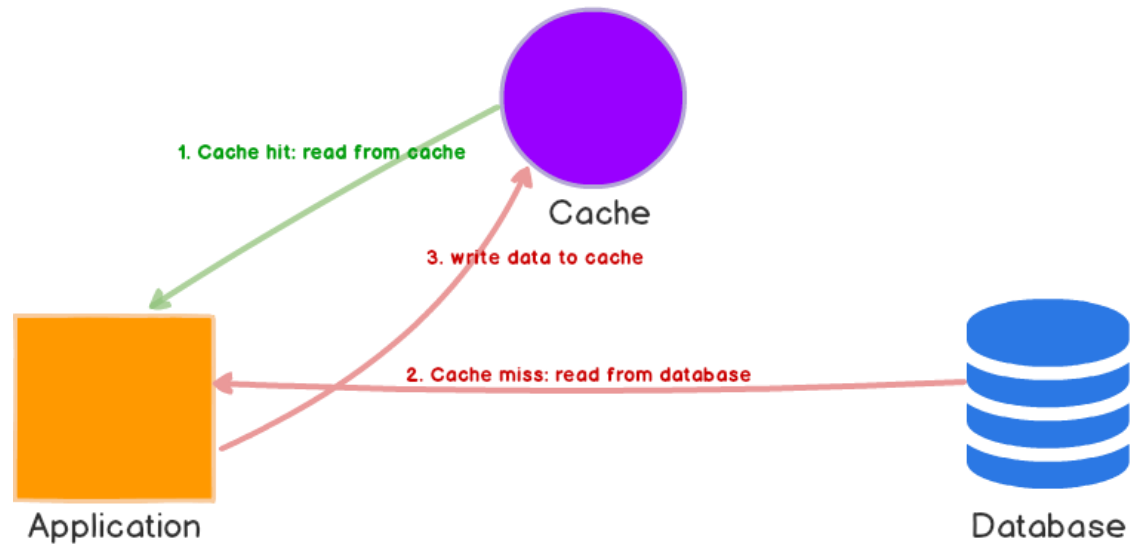
- First In First Out (FIFO)
- Last In First Out (LIFO)
- Least Recently Used (LRU)
- Most Recently Used (MRU)
- Least Frequently Used (LFU)
- Random Replacement (RR)

Cache update strategies

- Cache-aside (lazy loading)
- Read-through
- Write-through
- Write-behind (write-back)
- Write-around


Cache-aside (lazy loading)

Cache-Aside




Read-through

Read-Through




Write-through

Write-Through




Write-back

Write-Back



Asynchronism

Message Queue



The screenshot shows a product page for 'OSRAM Single Article for Release 1.4_06'. The product is a bottle of Coca-Cola. The page includes a 'FEEDBACK' button, a 'Beschikbaar: Online' status, and a 'Levering: Binnen 3 werkdagen' delivery time. The product has a rating of 4.7/5 stars based on 103 reviews. A 'Meer informatie' link is present. The page also features a 'Verantwoordingsvolter Genus' badge with '18 Jahren' and a 'Artikelnummer: 100023571'. A 'Hoe kan ik een beoordeling indienen?' (How can I submit a review?) dialog box is open at the bottom, explaining that reviews are automatically translated from the German Lidl Online Shop and that the user can leave a review in the Dutch Online Shop if they have one.

Scaling the DB


Replication: Master - slave

Cons:

- single point of failure
- Master slave cannot be scaled horizontally
- when failure, a slave has to be promoted to master to take over its place. No automatic failover
- Downtime and possibly loss of data when a master fails

Pros:

- You can split read and write requests to different servers



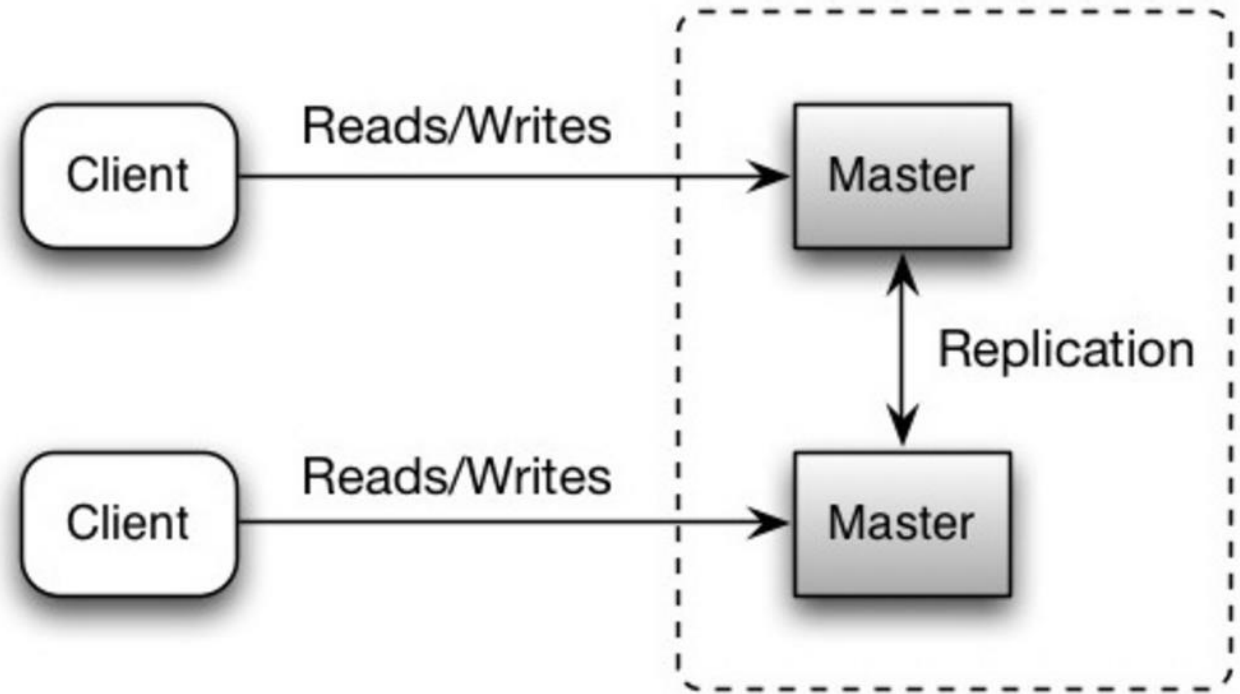
Replication: Multi – master

Cons:

- a load balancer or need of making changes to your application logic to determine where to write
- either loosely consistent or increased write latency due to synchronization

Pros:


- Failover semi-automatic because of multiple master nodes
- Distributes write load



Partitioning: Federation (Vertical)

Disadvantages

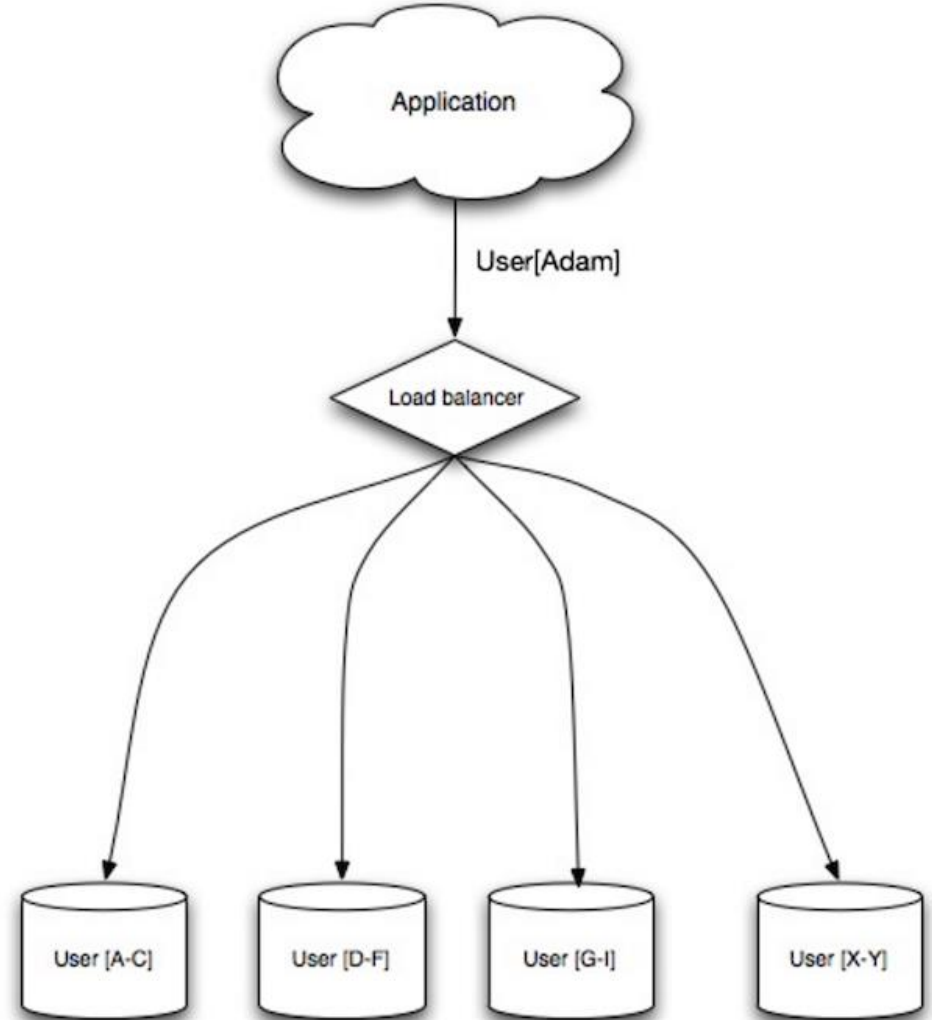
- not effective if your schema requires huge functions or tables
- update application logic to determine which database to read and write
- Joining data from two databases is more complex (FK)
- more hardware and additional complexity



Partitioning: Sharding (Horizontal)

Disadvantages:

- update application logic to work with shards; complex SQL queries.
- rebalancing
- joining data from multiple shards is more complex (FK)
- more hardware and additional complexity



25k signups in the first
day

scaling = replacing all
components of a car
while driving it at
100mph

moved db to its own
machine

but photos kept growing
and growing...

Instagram DB Scaling

vertical partitioning

photosdb > 60GB

horizontal partitioning!

aka: sharding

Instagram DB Scaling

what's painful about sharding?

1 data retrieval

2 what happens if one of your shards gets too big?


```
// 8 logical shards on 4 machines  
user_id % 8 = logical shard  
  
logical shards -> physical shard map  
{  
  0: A, 1: A,  
  2: C, 3: C,  
  4: B, 5: B,  
  6: D, 7: D  
}
```

can do this as long as you have more logical shards than physical ones

Instagram DB Scaling

Consistent Hashing

Hashring



Resources

- <https://github.com/donnemartin/system-design-primer> (Donne Martin github)
- Educative.io System Design course
- <https://www.slideshare.net/iammutex/scaling-instagram> (Scaling Instagram)
- <https://www.youtube.com/watch?v=zaRkONvyGr8> (Consistent Hashing)
- <https://svenbayer.blog/2018/09/30/accelerate-microservices-with-refresh-ahead-caching/>
- Google :)